

```
In [1]: %load_ext autoreload
        %autoreload 2
```

Final Evaluation

```
In [91]: import os
import sys
import json
import pickle
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import matplotlib.patches as patches
```

Load data

```
In [92]: ## here I'm loading predictions from pickled predictions
preds_path = '../data/predictions/predictions-final.pickle'
with open(preds_path, 'rb') as pickle_file:
    preds = pickle.load(pickle_file)
```

```
In [93]: ## load ground truths (val_gt.json)
with open('../data/predictions/val_gt.json') as json_file:
    val_gt = json.load(json_file)
```

```
In [94]: ## extract image ids
img_ids = {}
for vgt in val_gt['images']:
    img_ids[vgt['im_name']] = vgt['id']
```

fix the dictionary

```
cutoff = 1000
anno_dict2 = {}
for img in list(anno_dict.keys()):
    bboxes = anno_dict[img]

    bboxes2 = []
    for bbox in bboxes:
        if np.prod(bbox[2:]) > cutoff:
            bboxes2.append(bbox)
    if len(bboxes2) > 0:
        anno_dict2[img] = bboxes2

## get_annons
bboxes = []
for bb in anno_train[0, i][0][0][2]:
    if bb[0] > 0: # class_label = 1 means it is a person
        bboxes.append(bb[1:5]) # bbox format = [x, y, w, h]

## keep only images with persons
if bboxes != []:
    d[img_name] = bboxes
```

```
In [95]: val_gt['annotations'][0]
```

```
Out[95]: {'id': 1,
'image_id': 1,
'category_id': 1,
'iscrowd': 0,
'ignore': 0,
'bbox': [947, 406, 17, 40],
'vis_bbox': [950, 407, 14, 39],
'height': 40,
'vis_ratio': 0.802941176471}
```

```
In [96]: ## intersect with our predictions (441 images)
val_imgs = val_gt['images']
val_imgs = [val_img['im_name'] for val_img in val_imgs]

pred_imgs = list(preds.keys())
imgs = list(set(val_imgs) & set(pred_imgs))
len(imgs)
```

```
Out[96]: 418
```

```
In [247... ## subset val_gt and create test_gt
cutoff = 100
test_gt = {}
test_gt['categories'] = val_gt['categories']
test_gt['images'] = []
```

```

test_gt['annotations'] = []

for img in val_gt['images']:
    if img['im_name'] in imgs:
        test_gt['images'].append(img)

test_ids = []
for img in imgs:
    if img in list(img_ids.keys()):
        test_ids.append(img_ids[img])

for anno in val_gt['annotations']:
    if anno['image_id'] in test_ids:

        ## add area to anno (needed for default pycocotools eval)
        bbox = anno['bbox']
        area = int(bbox[2]) * int(bbox[3])
        anno['area'] = area

        ## cutoff small boxes
        if area > cutoff and anno['category_id'] > 0:
            test_gt['annotations'].append(anno)

```

```

In [248... ## check
test_imgs = test_gt['images']
test_imgs = [test_img['im_name'] for test_img in test_imgs]
test_imgs[:5], len(test_imgs)

```

```

Out[248... (['frankfurt_000000_000294_leftImg8bit.png',
'frankfurt_000000_000576_leftImg8bit.png',
'frankfurt_000000_001016_leftImg8bit.png',
'frankfurt_000000_001236_leftImg8bit.png',
'frankfurt_000000_001751_leftImg8bit.png'],
418)

```

```

In [249... test_gt['images'][0]

```

```

Out[249... {'id': 1,
'im_name': 'frankfurt_000000_000294_leftImg8bit.png',
'height': 1024,
'width': 2048}

```

```

In [250... ## GT bboxes
img_name = 'munster_000000_000019_leftImg8bit.png'
idx = test_imgs.index(img_name)
test1_img = test_gt['images'][idx]
img_name = test1_img['im_name']
img_id = test1_img['id']

bboxes = []
for anno in test_gt['annotations']:
    if anno['image_id'] == img_id and anno['category_id'] >= 1:
        bboxes.append(anno['bbox'])

```

```

In [251... # DT bboxes
bboxes_dt = np.round(preds[img_name]['boxes'])
bboxes_dt = bboxes_dt.tolist()

```

```

In [252... img = Image.open('../data/predictions/' + img_name)
plt.rcParams['figure.figsize'] = [12, 8]

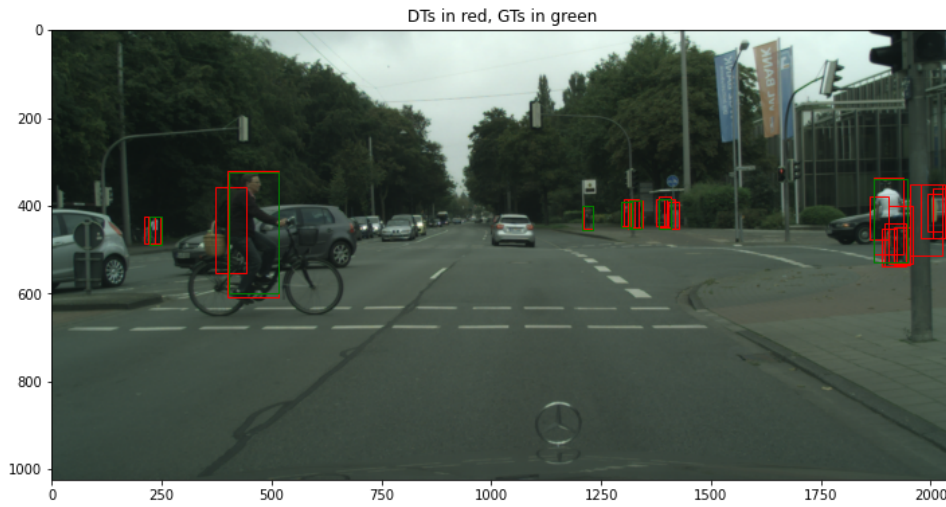
fig, ax = plt.subplots()
ax.imshow(img);

for bbox in bboxes_dt:
    x1, y1, x2, y2 = bbox
    w, h = x2 - x1, y2 - y1
    rect = patches.Rectangle(
        (x1, y1), w, h,
        linewidth=1, edgecolor='r', facecolor='none')
    ax.add_patch(rect)

# bbox = [x, y, w, h]
for bbox in bboxes:
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor='g', facecolor='none')
    ax.add_patch(rect)

plt.title('DTs in red, GTs in green')
plt.show()

```



```
In [253...  ## we can save the test_gt as json
with open('../data/predictions/test_gt.json', 'w', encoding='utf-8') as json_file:
    json.dump(test_gt, json_file, ensure_ascii=False, indent=4)
```

```
In [254...  ## see loadRes() and loadNumpyAnnotations() from COCO Class
## we need to provide [imageID, x1, y1, w, h, score, class] for each bbox:
test_dt = []
for img in imgs:
    bboxes = preds[img]['boxes']
    scores = preds[img]['scores']
    for i, bbox in enumerate(bboxes):
        x1, y1, x2, y2 = int(bbox[0]), int(bbox[1]), int(bbox[2]), int(bbox[3])
        w, h = x2 - x1, y2 - y1
        data = [img_ids[img], x1, y1, w, h, scores[i], 1]
        test_dt.append(data)
test_dt = np.array(test_dt)
```

```
In [255...  ## check
np.round(test_dt[0])
```

```
Out[255... array([3.060e+02, 1.416e+03, 3.520e+02, 9.100e+01, 2.180e+02, 1.000e+00,
        1.000e+00])
```

```
In [256...  ## check
# should be a lot more detections than gt bboxes
len(test_dt), len(test_gt['annotations'])
```

```
Out[256... (14773, 3601)
```

```
In [257...  print('About %.2f times more detected bboxes than there are gt bboxes'
      % (len(test_dt) / len(test_gt['annotations'])))
```

About 4.10 times more detected bboxes than there are gt bboxes

DEBUG

What I did:

- Converted CityPersons scripts in `eval_script` from Python2 to Python3 using `2to3 . -w . in ./eval_script`
- I compared the code and why did you divide `gt['vis_ratio']` by 100?

Evaluate

```
In [258...  ## Citypersons average miss rate measures
module_path = os.path.abspath(os.path.join('../src/eval_script/'))
if module_path not in sys.path:
    sys.path.append(module_path)

from coco import COCO
from eval_MR_multisetup import COCOeval

annType = 'bbox'
annFile = '../data/predictions/test_gt.json'
resFile = test_dt

res_file_path = '../data/predictions/results.txt'
res_file = open(res_file_path, 'w')

for id_setup in range(0, 4):
    cocoGt = COCO(annFile)
    cocoDt = cocoGt.loadRes(resFile)
    imgIds = sorted(cocoGt.getImgIds())
```

```

cocoEval = COCOeval(cocoGt, cocoDt, annType)
cocoEval.params.imgIds = imgIds
cocoEval.evaluate(id_setup)
cocoEval.accumulate()
cocoEval.summarize(id_setup, res_file)

res_file.close()

```

```

loading annotations into memory...
Done (t=0.01s)
creating index...
index created!
Loading and preparing results...
Converting ndarray to lists...
(14773, 7)
0/14773
DONE (t=0.06s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=1.21s).
Accumulating evaluation results...
DONE (t=0.00s).
Average Miss Rate (MR) @ Reasonable [ IoU=0.50 | height=[50:10000000000] | visibility=[0.65:10000000000.00]
] = 25.33%
loading annotations into memory...
Done (t=0.07s)
creating index...
index created!
Loading and preparing results...
Converting ndarray to lists...
(14773, 7)
0/14773
DONE (t=0.03s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=0.65s).
Accumulating evaluation results...
DONE (t=0.00s).
Average Miss Rate (MR) @ Reasonable_small [ IoU=0.50 | height=[50:75] | visibility=[0.65:10000000000.00] ] = 50.9
6%
loading annotations into memory...
Done (t=0.01s)
creating index...
index created!
Loading and preparing results...
Converting ndarray to lists...
(14773, 7)
0/14773
DONE (t=0.09s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=1.13s).
Accumulating evaluation results...
DONE (t=0.00s).
Average Miss Rate (MR) @ Reasonable_occ=heavy [ IoU=0.50 | height=[50:10000000000] | visibility=[0.20:0.65] ] = 63.
56%
loading annotations into memory...
Done (t=0.01s)
creating index...
index created!
Loading and preparing results...
Converting ndarray to lists...
(14773, 7)
0/14773
DONE (t=0.13s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=1.11s).
Accumulating evaluation results...
DONE (t=0.01s).
Average Miss Rate (MR) @ All [ IoU=0.50 | height=[20:10000000000] | visibility=[0.20:10000000000.00]
] = 41.96%

```

In [259...

```

## making the printout nicer..
print('Results: ')
res_file = open(res_file_path, 'r')
lines = res_file.readlines()
res_file.close()
lines = [line.replace('10000000000.00', 'inf') for line in lines]
lines = [line.replace('10000000000', 'inf') for line in lines]
lines = [line.strip() for line in lines]
for line in lines:
    new = ''
    for elt in line.split(' '):
        if elt:
            new += elt + ' '
    print(new)

```

```

Results:
Average Miss Rate (MR) @ Reasonable [ IoU=0.50 | height=[50:inf] | visibility=[0.65:inf] ] = 25.33%
Average Miss Rate (MR) @ Reasonable_small [ IoU=0.50 | height=[50:75] | visibility=[0.65:inf] ] = 50.96%
Average Miss Rate (MR) @ Reasonable_occ=heavy [ IoU=0.50 | height=[50:inf] | visibility=[0.20:0.65] ] = 63.56%
Average Miss Rate (MR) @ All [ IoU=0.50 | height=[20:inf] | visibility=[0.20:inf] ] = 41.96%

```

```
In [260... ## rewrite as a row to add it to the benchmark table
results = [line.split('=')[1] for line in lines]
results = [line.split('=')[1] for line in lines]
results.insert(0, ' x ')
results.insert(0, ' Our FasterRCNN ')
results = ['**' + result.strip() + '**') for result in results]
results = ' | '.join(results)
results = ' | ' + results + ' | '
results
```

```
Out[260... ' | **Our FasterRCNN** | *** | **25.33%** | **50.96%** | **63.56%** | **41.96%** | '
```

Benchmark

Method	External training data	MR (Reasonable)	MR (Reasonable_small)	MR (Reasonable_occ=heavy)	MR (All)
APD-pretrain	✓	7.31%	10.81%	28.07%	32.71%
Pedestron	✓	7.69%	9.16%	27.08%	28.33%
APD	x	8.27%	11.03%	35.45%	35.65%
YT-PedDet	x	8.41%	10.60%	37.88%	37.22%
STNet	x	8.92%	11.13%	34.31%	29.54%
MGAN	x	9.29%	11.38%	40.97%	38.86%
DVRNet	x	11.17%	15.62%	42.52%	40.99%
HBA-RCNN	x	11.26%	15.68%	39.54%	38.77%
OR-CNN	x	11.32%	14.19%	51.43%	40.19%
AdaptiveNMS	x	11.40%	13.64%	46.99%	38.89%
Reputation Loss	x	11.48%	15.67%	52.59%	39.17%
Cascade MS-CNN	x	11.62%	13.64%	47.14%	37.63%
Adapted FasterRCNN	x	12.97%	37.24%	50.47%	43.86%
MS-CNN	x	13.32%	15.86%	51.88%	39.94%
Our FasterRCNN	x	24.73%	47.35%	64.74%	52.72%

```
In [261... ## Test using default pycocotools measures
module_path = os.path.abspath(os.path.join('../src/pycocotools/'))
if module_path not in sys.path:
    sys.path.append(module_path)

from coco import COCO
from cocoEval import COCOeval

annType = 'bbox'
annFile = '../data/predictions/test_gt.json'
resFile = test_dt

cocoGt=COCO(annFile)
cocoDt = cocoGt.loadRes(resFile)

imgIds = sorted(cocoGt.getImgIds())
cocoEval = COCOeval(cocoGt, cocoDt, annType)

cocoEval.params.imgIds = imgIds
cocoEval.evaluate()

cocoEval.accumulate()
cocoEval.summarize()

loading annotations into memory...
Done (t=0.05s)
creating index...
index created!
Loading and preparing results...
Converting ndarray to lists...
(14773, 7)
0/14773
DONE (t=0.20s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=3.97s).
Accumulating evaluation results...
DONE (t=0.12s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.420
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.697
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.447
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.053
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.347
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.606
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.086
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.379
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.207
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.446
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.654
```

TODOS:

1. Why did we get worst results here compared to results on the cloud. Using `evaluate(model, data_loader_test, device=device)` when testing on the same set:

Update:

Average Precision	(AP)	@[IoU=0.50:0.95	area= all	maxDets=100]	= 0.461
Average Precision	(AP)	@[IoU=0.50	area= all	maxDets=100]	= 0.754
Average Precision	(AP)	@[IoU=0.75	area= all	maxDets=100]	= 0.492
Average Precision	(AP)	@[IoU=0.50:0.95	area= small	maxDets=100]	= 0.087
Average Precision	(AP)	@[IoU=0.50:0.95	area=medium	maxDets=100]	= 0.392
Average Precision	(AP)	@[IoU=0.50:0.95	area= large	maxDets=100]	= 0.616
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets= 1]	= 0.095
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets= 10]	= 0.417
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets=100]	= 0.550
Average Recall	(AR)	@[IoU=0.50:0.95	area= small	maxDets=100]	= 0.340
Average Recall	(AR)	@[IoU=0.50:0.95	area=medium	maxDets=100]	= 0.506
Average Recall	(AR)	@[IoU=0.50:0.95	area= large	maxDets=100]	= 0.663

In []: